


ML Principles & Practical Issues

The characteristics of good ML problems

- Clear use case




- Relevant data



- Decision making




The ML Pipeline



A Hypothetical case study

- A lab scientist wants to build an automated system that will allow her cat in and out of her office window and disallow dogs from entering through it



ML Problem formulation

- ML Problem Statement: Classify cats vs dogs correctly

Supervised Learning

$$y = f(x)$$

Predicted label

- cat
- dog

Input features from

- Images of cats and dogs

Machine learning function
that maps input to output

Build an ML classifier

- Collect data

What kind?

How much?

From which source?

Time and Expense/
Quality?

- Pre-process the image data

Resizing

Denoising


- Feature generation and data modelling

Handcrafted/ Statistical/
Deep learnt features

Which classifier?

- Evaluate the model

Which accuracy
metrics?



Practical issues



TRAIN



PREDICT

Predict
using trained
model



New Sample



Output

What could be the reasons?

Data

- Noisy / low quality data?
- Insufficient data volume?
- Poor data pre-processing?

Features


- Scaling required?
- Feature selection/
extraction required?

Model Training and Testing

- Insufficient training?
- Excessive training?
- Tuning of model parameters required?
- Algorithm needs to be changed?
- Testing method?

Holdout test set: The naïve approach


- Randomly split the entire dataset into:
 - **Training set:** A dataset used for training the model
 - **Test set (a.k.a validation set):** Data only used for testing the model



The three-way split

- **Training set**
 - A set of examples used for learning
- **Validation set**
 - A set of examples used to tune the parameters of a classifier
- **Test set**
 - A set of examples used only to assess the performance of a fully-trained classifier.

The three-way split



How to perform the split?

- How many examples in each data set?
 - **Training:** Typically 60-80% of data
 - **Test set:** Typically 20-30% of your data set
 - **Validation set:** Around 20% of data
- Examples
 - **3 way:** Training: 60%, Val: 20%, Test: 20%
 - **2 ways:** Training 70%, Test: 30%

Holdout summary


- **Positive**
 - Intuitive; Usually easy to perform; Considered the ideal method for evaluation
- **Drawbacks:**
 - In small datasets, you do not have the luxury of setting aside a portion of your data
 - The performance will be misleading if we had an unfortunate split

Common Splitting Strategies

- k-fold cross-validation




- Leave-one-out (n-fold cross validation)



Estimating model performance

K-fold cross validation



- The dataset is split into K partitions of equal size
- k-1 folds are used to train a model, and the holdout kth fold is used as the validation set
- This process is repeated and each of the folds is given an opportunity to be used as the holdout test set. A total of k models are fit and evaluated, and the performance of the model is calculated as the mean of these runs.



- Small datasets



- Computationally more expensive

Cross - Validation

- How do you summarize the performance? $E = \frac{1}{K} \sum_{i=1}^K E_i$
 - **Average:** Usually average of performance between experiments
- How many folds are needed?
 - **Common choice:** 5-fold or 10-fold cross-validation (Some nice numbers)
 - Large datasets → even 3-Fold cross-validation will do
 - Smaller datasets → bigger K. Why?
 - Leave-One-Out approach (K=n). K is equal to the number of examples n. Used for very small datasets

Thanks!!

Questions?